

Estudante: Clístenes Grizafis Bento

Professor: João Eugênio Marynowski

### **Atividade 1 (Salvem prints da tela para envio)**

- 1) `$ hadoop jar /usr/lib/hadoop-mapreduce/hadoop-mapreduce-examples.jar pi 1 100`
- 2) `$ hadoop jar /usr/lib/hadoop-mapreduce/hadoop-mapreduce-examples.jar pi 1 500000`
- 3) `$ hadoop jar /usr/lib/hadoop-mapreduce/hadoop-mapreduce-examples.jar pi 5 100000`

- Comentem sobre a diferença de execução dos exercícios 2 e 3  
(Quais diferenças? Tempo? Por que?)

Resposta:

Em relação ao pi 1 100 e os outros houve melhoria de precisão nos outros devido ao número de interações.

Comparando o tempo o pi 1 100 e o pi 1 500000 tiveram tempo aproximados, já o pi 1 500000 e o pi 5 100000 embora tenha a mesma quantidade de interações o segundo demorou um segundo mais devido ao fato de tempo para executar cinco processos paralelos de 100000 interações ao invés de um com 500000.

Parametros: pi 1 100

```
root@sandbox:~  
Using username "root".  
root@localhost's password:  
Last login: Sat Nov 19 06:07:25 2022  
[root@sandbox ~]# hadoop jar /usr/lib/hadoop-mapreduce/hadoop-mapreduce-examples.jar pi 1 100  
Number of Maps = 1  
Samples per Map = 100  
Wrote input for Map #0  
Starting Job  
22/11/19 06:09:35 INFO client.RMProxy: Connecting to ResourceManager at sandbox.hortonworks.com/10.0.2.15:8050  
22/11/19 06:09:36 INFO input.FileInputFormat: Total input paths to process : 1  
22/11/19 06:09:36 INFO mapreduce.JobSubmitter: number of splits:1  
22/11/19 06:09:36 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1668866818960_0001  
22/11/19 06:09:36 INFO impl.YarnClientImpl: Submitted application application_1668866818960_0001  
22/11/19 06:09:37 INFO mapreduce.Job: The url to track the job: http://sandbox.hortonworks.com:8088/proxy/application_1668866818960_0001/  
22/11/19 06:09:37 INFO mapreduce.Job: Running job: job_1668866818960_0001  
22/11/19 06:09:46 INFO mapreduce.Job: Job job_1668866818960_0001 running in uber mode : false  
22/11/19 06:09:46 INFO mapreduce.Job: map 0% reduce 0%  
22/11/19 06:09:52 INFO mapreduce.Job: map 100% reduce 0%  
22/11/19 06:09:59 INFO mapreduce.Job: map 100% reduce 100%  
22/11/19 06:10:00 INFO mapreduce.Job: Job job_1668866818960_0001 completed successfully  
22/11/19 06:10:00 INFO mapreduce.Job: Counters: 49  
File System Counters  
FILE: Number of bytes read=28  
FILE: Number of bytes written=199083  
FILE: Number of read operations=0  
FILE: Number of large read operations=0  
FILE: Number of write operations=0  
HDFS: Number of bytes read=277  
HDFS: Number of bytes written=215  
HDFS: Number of read operations=7  
HDFS: Number of large read operations=0  
HDFS: Number of write operations=3  
Job Counters  
Launched map tasks=1  
Launched reduce tasks=1  
Data-local map tasks=1  
Total time spent by all maps in occupied slots (ms)=4009  
Total time spent by all reduces in occupied slots (ms)=3895  
Total time spent by all map tasks (ms)=4009  
Total time spent by all reduce tasks (ms)=3895  
Total vcore-seconds taken by all map tasks=4009
```

```
root@sandbox:~  
HDFS: Number of write operations=3  
Job Counters  
  Launched map tasks=1  
  Launched reduce tasks=1  
  Data-local map tasks=1  
  Total time spent by all maps in occupied slots (ms)=4009  
  Total time spent by all reduces in occupied slots (ms)=3895  
  Total time spent by all map tasks (ms)=4009  
  Total time spent by all reduce tasks (ms)=3895  
  Total vcore-seconds taken by all map tasks=4009  
  Total vcore-seconds taken by all reduce tasks=3895  
  Total megabyte-seconds taken by all map tasks=1002250  
  Total megabyte-seconds taken by all reduce tasks=973750  
Map-Reduce Framework  
  Map input records=1  
  Map output records=2  
  Map output bytes=18  
  Map output materialized bytes=28  
  Input split bytes=159  
  Combine input records=0  
  Combine output records=0  
  Reduce input groups=2  
  Reduce shuffle bytes=28  
  Reduce input records=2  
  Reduce output records=0  
  Spilled Records=4  
  Shuffled Maps =1  
  Failed Shuffles=0  
  Merged Map outputs=1  
  GC time elapsed (ms)=144  
  CPU time spent (ms)=1570  
  Physical memory (bytes) snapshot=310878208  
  Virtual memory (bytes) snapshot=1830608896  
  Total committed heap usage (bytes)=155713536  
Shuffle Errors  
  BAD_ID=0  
  CONNECTION=0  
  IO_ERROR=0  
  WRONG_LENGTH=0  
  WRONG_MAP=0  
  WRONG_REDUCE=0  
File Input Format Counters  
  Bytes Read=118  
File Output Format Counters  
  Bytes Written=97  
Job Finished in 25.281 seconds  
Estimated value of Pi is 3.20000000000000000000  
[root@sandbox ~]#
```

Tempo: 25.281 segundos

Valor de PI: 3.20

Parametros: pi 1 500000

```
root@sandbox:~  
Using username "root".  
root@localhost's password:  
Last login: Sat Nov 19 06:07:25 2022  
[root@sandbox ~]# hadoop jar /usr/lib/hadoop-mapreduce/hadoop-mapreduce-examples.jar pi 1 100  
Number of Maps = 1  
Samples per Map = 100  
Wrote input for Map #0  
Starting Job  
22/11/19 06:09:35 INFO client.RMPProxy: Connecting to ResourceManager at sandbox.hortonworks.com/10.0.2.15:8050  
22/11/19 06:09:36 INFO input.FileInputFormat: Total input paths to process : 1  
22/11/19 06:09:36 INFO mapreduce.JobSubmitter: number of splits:1  
22/11/19 06:09:36 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1668866818960_0001  
22/11/19 06:09:36 INFO impl.YarnClientImpl: Submitted application application_1668866818960_0001  
22/11/19 06:09:37 INFO mapreduce.Job: The url to track the job: http://sandbox.hortonworks.com:8088/proxy/application_1668866818960_0001/  
22/11/19 06:09:37 INFO mapreduce.Job: Running job: job_1668866818960_0001  
22/11/19 06:09:46 INFO mapreduce.Job: Job job_1668866818960_0001 running in uber mode : false  
22/11/19 06:09:46 INFO mapreduce.Job: map 0% reduce 0%  
22/11/19 06:09:52 INFO mapreduce.Job: map 100% reduce 0%  
22/11/19 06:09:59 INFO mapreduce.Job: map 100% reduce 100%  
22/11/19 06:10:00 INFO mapreduce.Job: Job job_1668866818960_0001 completed successfully  
22/11/19 06:10:00 INFO mapreduce.Job: Counters: 49  
    File System Counters  
        FILE: Number of bytes read=28  
        FILE: Number of bytes written=199083  
        FILE: Number of read operations=0  
        FILE: Number of large read operations=0  
        FILE: Number of write operations=0  
        HDFS: Number of bytes read=277  
        HDFS: Number of bytes written=215  
        HDFS: Number of read operations=7  
        HDFS: Number of large read operations=0  
        HDFS: Number of write operations=3  
    Job Counters  
        Launched map tasks=1  
        Launched reduce tasks=1  
        Data-local map tasks=1  
        Total time spent by all maps in occupied slots (ms)=4009  
        Total time spent by all reduces in occupied slots (ms)=3895  
        Total time spent by all map tasks (ms)=4009
```

```
root@sandbox:~  
Data-local map tasks=1  
Total time spent by all maps in occupied slots (ms)=4421  
Total time spent by all reduces in occupied slots (ms)=4686  
Total time spent by all map tasks (ms)=4421  
Total time spent by all reduce tasks (ms)=4686  
Total vcore-seconds taken by all map tasks=4421  
Total vcore-seconds taken by all reduce tasks=4686  
Total megabyte-seconds taken by all map tasks=1105250  
Total megabyte-seconds taken by all reduce tasks=1171500  
Map-Reduce Framework  
Map input records=1  
Map output records=2  
Map output bytes=18  
Map output materialized bytes=28  
Input split bytes=159  
Combine input records=0  
Combine output records=0  
Reduce input groups=2  
Reduce shuffle bytes=28  
Reduce input records=2  
Reduce output records=0  
Spilled Records=4  
Shuffled Maps =1  
Failed Shuffles=0  
Merged Map outputs=1  
GC time elapsed (ms)=65  
CPU time spent (ms)=1770  
Physical memory (bytes) snapshot=340553728  
Virtual memory (bytes) snapshot=1782583296  
Total committed heap usage (bytes)=210239488  
Shuffle Errors  
BAD_ID=0  
CONNECTION=0  
IO_ERROR=0  
WRONG_LENGTH=0  
WRONG_MAP=0  
WRONG_REDUCE=0  
File Input Format Counters  
Bytes Read=118  
File Output Format Counters  
Bytes Written=97  
Job Finished in 24.504 seconds  
Estimated value of Pi is 3.14159200000000000000  
[root@sandbox ~]#  
[root@sandbox ~]#  
[root@sandbox ~]#
```

Tempo: 24.504segunndos

Valor de PI: 3.141592

Parametros: pi 5 100000

```
root@sandbox:~  
[root@sandbox ~]# hadoop jar /usr/lib/hadoop-mapreduce/hadoop-mapreduce-examples.jar pi 5 100000  
Number of Maps = 5  
Samples per Map = 100000  
  
Wrote input for Map #0  
Wrote input for Map #1  
Wrote input for Map #2  
Wrote input for Map #3  
Wrote input for Map #4  
Starting Job  
22/11/19 06:16:38 INFO client.RMProxy: Connecting to ResourceManager at sandbox.hortonworks.com/10.0.2.15:8050  
22/11/19 06:16:39 INFO input.FileInputFormat: Total input paths to process : 5  
22/11/19 06:16:39 INFO mapreduce.JobSubmitter: number of splits:5  
22/11/19 06:16:39 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1668866818960_0003  
22/11/19 06:16:39 INFO impl.YarnClientImpl: Submitted application application_1668866818960_0003  
22/11/19 06:16:39 INFO mapreduce.Job: The url to track the job: http://sandbox.hortonworks.com:8088/proxy/application_1668866818960_0003/  
22/11/19 06:16:39 INFO mapreduce.Job: Running job: job_1668866818960_0003  
22/11/19 06:16:48 INFO mapreduce.Job: Job job_1668866818960_0003 running in uber mode : false  
22/11/19 06:16:48 INFO mapreduce.Job: map 0% reduce 0%  
22/11/19 06:16:57 INFO mapreduce.Job: map 20% reduce 0%  
22/11/19 06:16:58 INFO mapreduce.Job: map 100% reduce 0%  
22/11/19 06:17:04 INFO mapreduce.Job: map 100% reduce 100%  
22/11/19 06:17:04 INFO mapreduce.Job: Job job_1668866818960_0003 completed successfully  
22/11/19 06:17:04 INFO mapreduce.Job: Counters: 49  
    File System Counters  
        FILE: Number of bytes read=116  
        FILE: Number of bytes written=597447  
        FILE: Number of read operations=0  
        FILE: Number of large read operations=0  
        FILE: Number of write operations=0  
        HDFS: Number of bytes read=1385  
        HDFS: Number of bytes written=215  
        HDFS: Number of read operations=23  
        HDFS: Number of large read operations=0  
        HDFS: Number of write operations=3  
    Job Counters  
        Launched map tasks=5  
        Launched reduce tasks=1  
        Data-local map tasks=5
```

```
root@sandbox:~  
Launched reduce tasks=1  
Data-local map tasks=5  
Total time spent by all maps in occupied slots (ms)=38259  
Total time spent by all reduces in occupied slots (ms)=3869  
Total time spent by all map tasks (ms)=38259  
Total time spent by all reduce tasks (ms)=3869  
Total vcore-seconds taken by all map tasks=38259  
Total vcore-seconds taken by all reduce tasks=3869  
Total megabyte-seconds taken by all map tasks=9564750  
Total megabyte-seconds taken by all reduce tasks=967250  
Map-Reduce Framework  
Map input records=5  
Map output records=10  
Map output bytes=90  
Map output materialized bytes=140  
Input split bytes=795  
Combine input records=0  
Combine output records=0  
Reduce input groups=2  
Reduce shuffle bytes=140  
Reduce input records=10  
Reduce output records=0  
Spilled Records=20  
Shuffled Maps =5  
Failed Shuffles=0  
Merged Map outputs=5  
GC time elapsed (ms)=392  
CPU time spent (ms)=4010  
Physical memory (bytes) snapshot=1123315712  
Virtual memory (bytes) snapshot=5298356224  
Total committed heap usage (bytes)=628097024  
Shuffle Errors  
BAD_ID=0  
CONNECTION=0  
IO_ERROR=0  
WRONG_LENGTH=0  
WRONG_MAP=0  
WRONG_REDUCE=0  
File Input Format Counters  
Bytes Read=590  
File Output Format Counters  
Bytes Written=97  
Job Finished in 26.289 seconds  
Estimated value of Pi is 3.1415920000000000000000  
[root@sandbox ~]#  
[root@sandbox ~]#
```

Tempo: 26.289 segundos

Valor de PI: 3.141592